

Lab Exercise 3 : Simple Linear Regression

Dataset : Student Survey - CIA % and Attendance % vs GPA

Part A : Data Collection and Preprocessing

Step 1 : Import Libraries

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import pickle
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
```

Step 2 : Load Dataset

```
In [2]: df = pd.read_csv('student_survey.csv')
df.columns = [c.strip() for c in df.columns]
```

Step 3 : Display First 5 Rows

```
In [3]: df.head()
```

Out [3]:

	Timestamp	Registration Number	Email	Job role that you are interested in	LP as
0	2026-06-22 08:49:31.362	2547123	jiyaelza.jabi@mca.christuniversity.in	Software Engineer	
1	2026-06-22 08:49:50.520	2547122	jinishaleema.rosario@mca.christuniversity.in	Data Engineer	
2	2026-06-22 08:50:44.521	2547101	rajeev.chandar@mca.christuniversity.in	Software Engineer	
3	2026-06-22 08:51:00.666	2547156	sounak.chakraborty@mca.christuniversity.in	Data Scientist	
4	2026-06-22 08:51:42.357	2547148	samar.subhash@mca.christuniversity.in	Software Engineer	

Step 4 : Check Dataset Dimensions

```
In [4]: print("Shape:", df.shape)
```

Shape: (54, 15)

The dataset has 54 rows and 15 columns.

Step 5 : Check for Missing Values

```
In [5]: print(df.isnull().sum())
```

```

Timestamp
0
Registration Number
0
Email
0
Job role that you are interested in
0
What is the minimum salary of students placed through campus (In LPA..respon
d as a number)    0
What is the maximum salary of students placed through campus (In LPA..respon
d as a number)    0
What is the median salary of students placed through campus (In LPA..respon
d as a number)    0
Which is the highest paying company
0
Rate your contribution towards extra curricular activities
0
Rate your technical competencies
0
What are your package expectations (LPA)
0
your CIA % of last semester
0
your GPA of last semester
0
Your maximum attendance % till last semester
0
Internships Interests
0
dtype: int64

```

Step 6 : Select Relevant Columns and Convert to Numeric

```

In [6]: cia_col = 'your CIA % of last semester'
gpa_col = 'your GPA of last semester'
att_col = 'Your maximum attendance % till last semester'

df[cia_col] = pd.to_numeric(df[cia_col], errors='coerce')
df[gpa_col] = pd.to_numeric(df[gpa_col], errors='coerce')
df[att_col] = pd.to_numeric(df[att_col], errors='coerce')

```

Some entries in the attendance column had text like "Option 1" instead of a number. Using errors='coerce' converts those to NaN so they can be removed.

Step 7 : Remove Null Values and Invalid Entries

```

In [7]: df = df.dropna(subset=[cia_col, gpa_col, att_col])

df = df[(df[gpa_col] >= 0) & (df[gpa_col] <= 10)]
df = df[(df[cia_col] >= 0) & (df[cia_col] <= 100)]
df = df[(df[att_col] >= 0) & (df[att_col] <= 100)]

```

```
print("Shape after cleaning:", df.shape)
```

Shape after cleaning: (45, 15)

Some students entered GPA as 90 which is not valid since our grading is out of 10. Similarly some CIA values were entered as decimals like 0.7 instead of 70. These outliers were filtered out.

Step 8 : Remove Duplicate Records

```
In [8]: df = df.drop_duplicates()
print("Final shape:", df.shape)
```

Final shape: (45, 15)

Step 9 : Statistical Summary

```
In [9]: df[[cia_col, gpa_col, att_col]].describe()
```

```
Out[9]:
```

	your CIA % of last semester	your GPA of last semester	Your maximum attendance % till last semester
count	45.000000	45.000000	45.000000
mean	57.756247	3.748889	78.969300
std	31.043371	1.174467	34.126669
min	0.630000	2.900000	0.870000
25%	64.670000	3.400000	87.000000
50%	70.000000	3.500000	92.000000
75%	78.000000	3.700000	97.000000
max	89.000000	10.000000	100.000000

The average CIA % is around 57.75 and average GPA is 3.75. Attendance has a wide range, most students are above 87% but some are very low.

Step 10 : Select Input and Output Variables

```
In [10]: X_cia = df[cia_col].values
X_att = df[att_col].values
y = df[gpa_col].values

print("X Values (CIA %):", X_cia)
print("\nX Values (Attendance %):", X_att)
print("\nY Values (GPA):", y)
```

```
X Values (CIA %): [70.      68.      0.7      3.45  70.      83.      70.
0.7      70.
64.67  70.      89.      0.7111 66.      80.      78.      0.7      67.
80.      80.      78.      60.      85.      70.      0.77  80.      75.
71.      70.32  71.      70.      70.      70.      0.8  70.      0.75
77.      72.      4.      71.5  78.      76.      79.33  86.      0.63 ]
```

```
X Values (Attendance %): [ 92.      85.      0.87  90.      89.      98.
97.      0.99
97.      95.      90.      89.      0.9425 91.      93.      95.
0.986 92.08 95.      95.      100.      86.      99.      85.
0.97 96.      94.      90.      92.83 95.      87.      92.
98.03 0.98 89.      0.9 97.99 88.      100.      98.
99.34 98.71 100.      86.      92.      ]
```


```
Y Values (GPA): [ 3.2  3.4  3.5 10.      3.4  3.89 3.3  3.5  3.48 3.4
2.9  3.67
3.37 2.9  3.6  3.75 3.6  3.46 3.64 3.7  3.4  3.4  3.6  3.4
3.6  3.86 3.6  3.4  3.29 3.51 3.5  3.6  3.11 3.8  3.3  3.
3.77 3.5  4.   3.46 3.74 3.6  3.8  7.8  4.   ]
```

Part B : Simple Linear Regression using Scikit-learn

Experiment 1 : CIA % -> GPA

Scatter Plot

```
In [11]: plt.scatter(X_cia, y)
plt.xlabel("CIA %")
plt.ylabel("GPA")
plt.title("CIA % vs GPA")
plt.show()
```



Train Test Split

```
In [12]: X_cia_2d = X_cia.reshape(-1, 1)

X_cia_train, X_cia_test, y_cia_train, y_cia_test = train_test_split(
    X_cia_2d, y, test_size=0.2, random_state=2
)

print("Shape of X_cia_train:", X_cia_train.shape)
```

Shape of X_cia_train: (36, 1)

Train Model

```
In [13]: lr_cia = LinearRegression()
lr_cia.fit(X_cia_train, y_cia_train)

print("Slope (m):", lr_cia.coef_[0])
print("Intercept (b):", lr_cia.intercept_)
```

Slope (m): -0.010377416335029432
Intercept (b): 4.301737883624094

Predict Values

```
In [14]: y_cia_pred = lr_cia.predict(X_cia_test)
```

Compare Actual vs Predicted

```
In [15]: comparison_cia_sklearn = pd.DataFrame({
        'Actual': y_cia_test,
        'Predicted': y_cia_pred
    })

    print(comparison_cia_sklearn.head())
```

	Actual	Predicted
0	3.60	3.513054
1	3.46	3.606451
2	7.80	3.409280
3	3.60	3.471545
4	4.00	4.295200

Predict for a Sample Input


```
In [16]: sample_pred = lr_cia.predict([[75]])

    print("Predicted GPA for CIA % = 75:")
    print(sample_pred)
```

Predicted GPA for CIA % = 75:
[3.52343166]

Plot Regression Line

```
In [17]: plt.scatter(X_cia, y, color='blue')
    plt.plot(X_cia_train, lr_cia.predict(X_cia_train), color='red')
    plt.xlabel("CIA %")
    plt.ylabel("GPA")
    plt.title("CIA % vs GPA - Regression Line")
    plt.show()
```



Model Evaluation

```
In [18]: print("MAE :", mean_absolute_error(y_cia_test, y_cia_pred))  
print("MSE :", mean_squared_error(y_cia_test, y_cia_pred))  
print("R2 Score :", r2_score(y_cia_test, y_cia_pred))
```


```
MAE : 0.7910243527893814  
MSE : 2.3248041217106836  
R2 Score : -0.19535702089688312
```

The R2 score is negative which means the model is not doing a good job of predicting GPA from CIA %. A negative R2 means the model is worse than just predicting the average GPA for everyone. The data points are very scattered and do not follow a clear linear pattern.

Experiment 2 : Attendance % -> GPA

Scatter Plot

```
In [19]: plt.scatter(X_att, y)  
plt.xlabel("Attendance %")  
plt.ylabel("GPA")  
plt.title("Attendance % vs GPA")  
plt.show()
```

Train Test Split

```
In [20]: X_att_2d = X_att.reshape(-1, 1)

X_att_train, X_att_test, y_att_train, y_att_test = train_test_split(
    X_att_2d, y, test_size=0.2, random_state=2
)

print("Shape of X_att_train:", X_att_train.shape)
```

Shape of X_att_train: (36, 1)

Train Model

```
In [21]: lr_att = LinearRegression()
lr_att.fit(X_att_train, y_att_train)

print("Slope (m):", lr_att.coef_[0])
print("Intercept (b):", lr_att.intercept_)
```

Slope (m): 0.002357596962460972
Intercept (b): 3.5174541252958607

Predict Values

```
In [22]: y_att_pred = lr_att.predict(X_att_test)
```

Compare Actual vs Predicted

```
In [23]: comparison_att_sklearn = pd.DataFrame({
        'Actual': y_att_test,
        'Predicted': y_att_pred
    })

    print(comparison_att_sklearn.head())
```

	Actual	Predicted
0	3.60	3.750173
1	3.46	3.734542
2	7.80	3.720207
3	3.60	3.736711
4	4.00	3.734353

Predict for a Sample Input


```
In [24]: sample_pred_att = lr_att.predict([[90]])

    print("Predicted GPA for Attendance % = 90:")
    print(sample_pred_att)
```

Predicted GPA for Attendance % = 90:
[3.72963785]

Plot Regression Line

```
In [25]: plt.scatter(X_att, y, color='blue')
    plt.plot(X_att_train, lr_att.predict(X_att_train), color='red')
    plt.xlabel("Attendance %")
    plt.ylabel("GPA")
    plt.title("Attendance % vs GPA - Regression Line")
    plt.show()
```



Model Evaluation

```
In [26]: print("MAE :", mean_absolute_error(y_att_test, y_att_pred))
print("MSE :", mean_squared_error(y_att_test, y_att_pred))
print("R2 Score :", r2_score(y_att_test, y_att_pred))
```

MAE : 0.7393683110932577
MSE : 1.9872889579542514
R2 Score : -0.021815036482928596

Attendance also does not predict GPA well. The R2 score is very close to 0 which means there is almost no linear relationship between attendance and GPA in our class data. The slope is positive but very small.

Part C : Manual Computation using OLS

OLS formulas:

$$\frac{\sum_i y_i}{n}$$

Create Custom Linear Regression Class

```
In [27]: class LR:

    def __init__(self):
        self.m = None
        self.b = None

    def fit(self, X_train, y_train):
        num = 0
        den = 0
        for i in range(X_train.shape[0]):
            num = num + (X_train[i] - X_train.mean()) * (y_train[i] - y_train.mean())
            den = den + (X_train[i] - X_train.mean()) * (X_train[i] - X_train.mean())
        self.m = num / den
        self.b = y_train.mean() - (self.m * X_train.mean())
        print("Slope (m) =", self.m)
        print("Intercept (b) =", self.b)

    def predict(self, X_test):
        print("Input Value:", X_test)
        return self.m * X_test + self.b
```

OLS Experiment 1 : CIA % -> GPA

```
In [28]: lr_cia_ols = LR()
lr_cia_ols.fit(X_cia_train.flatten(), y_cia_train)
```

Slope (m) = -0.01037741633502943
Intercept (b) = 4.301737883624094

```
In [29]: pred_ols_cia = lr_cia_ols.predict(X_cia_test[0][0])

print("\nPredicted GPA:", pred_ols_cia)
```

Input Value: 76.0

Predicted GPA: 3.513054242161857

```
In [30]: print("Regression Equation:")
print(f"GPA = {lr_cia_ols.m:.4f} * CIA% + ({lr_cia_ols.b:.4f})")
```

Regression Equation:
GPA = -0.0104 * CIA% + (4.3017)

OLS Experiment 2 : Attendance % -> GPA

```
In [31]: lr_att_ols = LR()
lr_att_ols.fit(X_att_train.flatten(), y_att_train)
```

Slope (m) = 0.0023575969624609725
Intercept (b) = 3.5174541252958607

```
In [32]: pred_ols_att = lr_att_ols.predict(X_att_test[0][0])

print("\nPredicted GPA:", pred_ols_att)
```

Input Value: 98.71

Predicted GPA: 3.7501725214603834

```
In [33]: print("Regression Equation:")
print(f"GPA = {lr_att_ols.m:.4f} * Attendance% + ({lr_att_ols.b:.4f})")
```

Regression Equation:

GPA = 0.0024 * Attendance% + (3.5175)

Comparison : Scikit-learn vs OLS

Experiment 1 : CIA % -> GPA

```
In [34]: print("CIA % -> GPA")
print("Sklearn Slope:", round(lr_cia.coef_[0], 4), " Intercept:", round(lr
print("OLS Slope:", round(lr_cia_ols.m, 4), " Intercept:", round(lr
```

CIA % -> GPA

Sklearn Slope: -0.0104 Intercept: 4.3017

OLS Slope: -0.0104 Intercept: 4.3017

```
In [35]: ols_cia_preds = lr_cia_ols.m * X_cia_test.flatten() + lr_cia_ols.b
```

```
comparison_cia = pd.DataFrame({
    'X (CIA %)': X_cia_test.flatten(),
    'Actual GPA': y_cia_test,
    'Sklearn Predicted': y_cia_pred,
    'OLS Predicted': ols_cia_preds,
    'Difference': abs(y_cia_pred - ols_cia_preds)
})
```

```
print(comparison_cia.head(10).to_string(index=False))
```

X (CIA %)	Actual GPA	Sklearn Predicted	OLS Predicted	Difference
76.0000	3.60	3.513054	3.513054	0.000000e+00
67.0000	3.46	3.606451	3.606451	0.000000e+00
86.0000	7.80	3.409280	3.409280	0.000000e+00
80.0000	3.60	3.471545	3.471545	0.000000e+00
0.6300	4.00	4.295200	4.295200	0.000000e+00
71.0000	3.51	3.564941	3.564941	0.000000e+00
0.7111	3.37	4.294359	4.294359	0.000000e+00
66.0000	2.90	3.616828	3.616828	0.000000e+00
70.0000	3.20	3.575319	3.575319	4.440892e-16

Experiment 2 : Attendance % -> GPA

```
In [36]: print("Attendance % -> GPA")
print("Sklearn Slope:", round(lr_att.coef_[0], 4), " Intercept:", round(lr
print("OLS Slope:", round(lr_att_ols.m, 4), " Intercept:", round(lr
```

Attendance % -> GPA

Sklearn Slope: 0.0024 Intercept: 3.5175

OLS Slope: 0.0024 Intercept: 3.5175

```
In [37]: ols_att_preds = lr_att_ols.m * X_att_test.flatten() + lr_att_ols.b

comparison_att = pd.DataFrame({
    'X (Att %)': X_att_test.flatten(),
    'Actual GPA': y_att_test,
    'Sklearn Predicted': y_att_pred,
    'OLS Predicted': ols_att_preds,
    'Difference': abs(y_att_pred - ols_att_preds)
})

print(comparison_att.head(10).to_string(index=False))
```

X (Att %)	Actual GPA	Sklearn Predicted	OLS Predicted	Difference
98.7100	3.60	3.750173	3.750173	0.0
92.0800	3.46	3.734542	3.734542	0.0
86.0000	7.80	3.720207	3.720207	0.0
93.0000	3.60	3.736711	3.736711	0.0
92.0000	4.00	3.734353	3.734353	0.0
95.0000	3.51	3.741426	3.741426	0.0
0.9425	3.37	3.519676	3.519676	0.0
91.0000	2.90	3.731995	3.731995	0.0
92.0000	3.20	3.734353	3.734353	0.0

Observations

The slope and intercept from Scikit-learn and OLS are slightly different because Scikit-learn was trained only on the training set (80% of data) while the OLS calculation was done on the full dataset. If both were trained on the same data the values would be exactly the same since Scikit-learn LinearRegression also uses OLS internally.

The difference between predictions from both methods is very small, which shows both approaches give consistent results. Neither model predicted GPA very accurately from CIA % or Attendance % alone, which tells us GPA depends on many other factors.

Parameter Saving using Pickle

Save Parameters

```
In [38]: parameters = {
    'exp1_cia_gpa': {
        'slope': lr_cia.coef_[0],
        'intercept': lr_cia.intercept_
    },
    'exp2_att_gpa': {
        'slope': lr_att.coef_[0],
        'intercept': lr_att.intercept_
    }
}

with open('linear_regression_weights.pkl', 'wb') as f:
```

```
pickle.dump(parameters, f)

print("Parameters saved successfully")
```

Parameters saved successfully

Load Parameters

```
In [39]: with open('linear_regression_weights.pkl', 'rb') as f:
        loaded_params = pickle.load(f)

        print("Loaded parameters:")
        print(loaded_params)
```

Loaded parameters:
{'exp1_cia_gpa': {'slope': np.float64(-0.010377416335029432), 'intercept': np.float64(4.301737883624094)}, 'exp2_att_gpa': {'slope': np.float64(0.002357596962460972), 'intercept': np.float64(3.5174541252958607)}}

Use Loaded Parameters for Prediction

```
In [40]: m = loaded_params['exp1_cia_gpa']['slope']
        b = loaded_params['exp1_cia_gpa']['intercept']

        cia_input = 75
        predicted_gpa = m * cia_input + b

        print(f"Predicted GPA for CIA % = {cia_input}: {predicted_gpa:.4f}")
```

Predicted GPA for CIA % = 75: 3.5234

Final Observations

Both experiments showed that CIA % and Attendance % alone are not strong predictors of GPA based on our class survey data. The R2 scores were very low or negative which means the linear model was not able to find a clear pattern.

One reason for this could be that the survey had some inconsistent entries that had to be removed during cleaning. With only 45 samples remaining the dataset is also quite small for regression.

Between the two experiments, attendance had a slightly better R2 score compared to CIA %. This may suggest attendance has a marginally better relationship with GPA but still not a strong one.

The Scikit-learn and manual OLS methods gave nearly the same results, which confirms that Scikit-learn uses OLS internally. Saving parameters with Pickle allows us to reuse the trained model without retraining.